

▼ LEARNING OF SEMANTICALLY EQUIVALENT QUESTIONS

```
# importing libraries
import keras
import nltk
import tensorflow as tf
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
import matplotlib.pyplot as plt
```

☞ Using TensorFlow backend.

```
# import dataset from drive
from google.colab import drive
drive.mount('/content/gdrive')
```

☞ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=

Enter your authorization code:

.....

Mounted at /content/gdrive

▼ IMPORTING DATASET FROM KAGGLE (optional)

The dataset can be downloaded from this website: <https://www.kaggle.com/c/quora-question-pairs/data>

The below cells shows how the dataset can be extracted to the notebook from kaggle. To import the dataset from kaggle.json from Kaggle -- MyAccount -- Create New API Token - auto downloads as "kaggle.json"

```
from google.colab import files
```

```
files.upload() #this will prompt you to upload the kaggle.json to the notebook
```

☞ kaggle.json

- **kaggle.json**(application/json) - 66 bytes, last modified: 7/25/2019 - 100% done

Saving kaggle.json to kaggle.json

```
{'kaggle.json': b'{"username":"hshreedhar","key":"e96183fe90a76e704596438aec7c338"}
```

```
!pip install -q kaggle
!mkdir -p ~/.kaggle
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

☞ kaggle.json

```
import os
os.chdir('/content/gdrive/My Drive/Kaggle/kaggle-questionpair/input')
```

```
!kaggle competitions download -c quora-question-pairs --force
```

```

[ ] Downloading sample_submission.csv.zip to /content/gdrive/My Drive/Kaggle/kaggle-q
 20% 1.00M/4.95M [00:00<00:01, 2.20MB/s]
100% 4.95M/4.95M [00:00<00:00, 10.1MB/s]
Downloading train.csv.zip to /content/gdrive/My Drive/Kaggle/kaggle-questionpair/
 80% 17.0M/21.2M [00:01<00:04, 1.05MB/s]
100% 21.2M/21.2M [00:01<00:00, 14.2MB/s]
Downloading test.csv.zip to /content/gdrive/My Drive/Kaggle/kaggle-questionpair/i
 93% 105M/112M [00:08<00:00, 10.1MB/s]
100% 112M/112M [00:08<00:00, 13.7MB/s]

```

▼ PREPARING THE DATASET

```

# Below is a sample dataset
# Dataset for training and validation
df = pd.read_csv('../input/train.csv') # read the CSV file
df.head() # to vieww the first 5 contents in a file

```

```

[ ]

```

	id	qid1	qid2	question1	question2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guid
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian
2	2	5	6	How can I increase the speed of my internet co...	How can Internet speed be increa
3	3	7	8	Why am I mentally very lonely? How can I solve...	Find the remainder when [math]2
4	4	9	10	Which one dissolve in water quikly sugar, salt...	Which fish would surv

We are setting the size of the training dataset in the block below. NUM_OF_SAMPLES can be changed as desired. NUM_OF_SAMPLES is the number of data points from each class. The total number of datapoints is 2*NUM_OF_SAI

```

# to get equal number of duplicate and non duplicate question pairs in the dataset
NUM_OF_SAMPLES = 10000 # Using smaller sample for training set, for each subset.

df_dup = df[df['is_duplicate'] == 1]
df_dup = df_dup[0:NUM_OF_SAMPLES]

df_diff = df[df['is_duplicate'] == 0]
df_diff = df_diff[0:NUM_OF_SAMPLES]

len(df_dup)

```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

▼ Limiting the size of the data set for an unbiased training

in the dataset, 60% of the question-pairs do not have a duplicate pair and 40% have a duplicate pair. We are creating qustion-pair with and without duplicates for an unbiased training. Off this sub-dataset that is created, we'll use 80% f

```
# here, we are spitting the dataset in training and validation, with 80% of the question-pai

num_of_train = int(round(0.8*NUM_OF_SAMPLES))      # Number of samples for training
num_of_valid = int(NUM_OF_SAMPLES - num_of_train)  # Number of samples for validation

# Choosing NUM_OF_SAMPLES random question pair for both duplicate and non-duplicate labels
dup_indices = np.random.choice(df_dup.index,size=num_of_train)
df_dup_train = df_dup.loc[dup_indices]

rem_indices = set(df_dup.index).difference(set(dup_indices)) # Remaining indices
dup_valid_indices = np.random.choice(list(rem_indices),num_of_valid)
df_dup_valid = df_dup.loc[dup_valid_indices]

diff_indices = np.random.choice(df_diff.index,size=num_of_train)
df_diff_train = df_diff.loc[diff_indices]

rem_indices = set(df_diff.index).difference(set(diff_indices)) # Remaining indices
diff_valid_indices = np.random.choice(list(rem_indices),num_of_valid)
df_diff_valid = df_diff.loc[diff_valid_indices]

# Selecting Training and Validation data points
df_train = pd.concat([df_dup_train,df_diff_train])
df_valid = pd.concat([df_dup_valid, df_diff_valid])

df_train_label = df_dup_train['is_duplicate']
df_train_label = df_train_label.append(df_diff_train['is_duplicate'])

df_valid_label = df_dup_valid['is_duplicate']
df_valid_label = df_valid_label.append(df_diff_valid['is_duplicate'])

df_train = df_train[['question1','question2']]
df_valid = df_valid[['question1', 'question2']]
```

Here, we have our dataset ready for training "df_train" and "df_valid" for validation. Duplicate values are the labels

```
# this block is to fetch the total number of duplicate question pair in the training set
print('Total number of question pair in the training set: {}'.format(len(df_train)))
```

```
➞ Total number of question pair in the training set: 16000
```

```
print(df_train.shape)
print(df_valid.shape)
```

```
➞ (16000, 2)
    (4000, 2)
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

```
# cleaning up of the dataset
import re
import string
```

```
re_punc = re.compile('[%s]' % re.escape(string.punctuation))
re_print = re.compile('[^%s]' % re.escape(string.printable))
```

```
def pre_process_line (line):
    # tokenize on white space
    try:
        line = line.split()
        # convert to lowercase
        line = [word.lower() for word in line]
        # remove punctuation from each token
        line = [re_punc.sub('', w) for w in line]
        # remove non-printable chars from each token
        line = [re_print.sub('', w) for w in line]
        # remove tokens with numbers in them
        line = [word for word in line if word.isalpha()]
        return ' '.join(line)
    except:
        return ' '

def pre_process_questions(ques):
    df_processed = ques.copy(deep=True)
    for i in range(ques.shape[0]):
        line = pre_process_line(ques.iloc[i])
        df_processed.iloc[i] = line
    return df_processed

df_train_q1 = df_train['question1'].apply(pre_process_line)
df_train_q2 = df_train['question2'].apply(pre_process_line)
df_valid_q1 = df_valid['question1'].apply(pre_process_line)
df_valid_q2 = df_valid['question2'].apply(pre_process_line)
```

▼ Encoding the questions

The word embedding (from q1 and q2) are split to training and validation datasets. Further, the words within the que

```
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences

all_qs = df_train_q1.tolist() + df_train_q2.tolist() + df_valid_q1.tolist() + df_valid_q2.to

tokenizer = Tokenizer()
tokenizer.fit_on_texts(all_qs)
vocab_size = len(tokenizer.word_index) + 1
encoded_qs = tokenizer.texts_to_sequences(all_qs)

MAX_QUES_LEN = 300
padded_qs = pad_sequences(encoded_qs, maxlen=MAX_QUES_LEN, padding='post')

# Divide the questions

encoded_train_q1 = padded_qs[0:(2*num_of_train)]
encoded_train_q2 = padded_qs[(2*num_of_train):(4*num_of_train)]
encoded_valid_q1 = padded_qs[(4*num_of_train):(4*num_of_train+2*num_of_valid)]
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

```
encoded_valid_q1.shape
```

```
↳ (4000, 300)
```

▼ Initializing Embedding layer using GloVe

```
import gensim
from gensim.test.utils import datapath, get_tmpfile
from gensim.models import KeyedVectors
from gensim.scripts.glove2word2vec import glove2word2vec

glove_file = datapath('/content/gdrive/My Drive/NLP- Deep Learning/glove.6B.300d.txt')
word2vec_glove_file = get_tmpfile("glove.6B.300d.word2vec.txt")
glove2word2vec(glove_file, word2vec_glove_file)
WORD2VEC = KeyedVectors.load_word2vec_format(word2vec_glove_file)
```

```
↳ /usr/local/lib/python3.6/dist-packages/smart_open/smart_open_lib.py:398: UserWarning
'See the migration notes for details: %s' % _MIGRATION_NOTES_URL
```

```
EMBEDDING_DIM = 300 # embedding dimentions are set to 300 as we are training on a 300 diment

def embeddings_init(shape, dtype=None):
    embedding_matrix = np.zeros((vocab_size, EMBEDDING_DIM))
    not_present = []
    for word, idx in tokenizer.word_index.items():
        try:
            vec = WORD2VEC[word]
            if vec is not None:
                embedding_matrix[idx] = vec
        except:
            not_present.append(word)
            continue
    print(len(not_present))
    return embedding_matrix
```

▼ TRAIN THE MODEL

```
import keras
from keras import backend as K
from keras.layers import Input, Dense, Conv1D, MaxPooling1D, Flatten, Embedding
from keras import models
from keras.optimizers import SGD
```

```
def get_siamese_model(input_shape):
    """
        Model architecture
    """
```

```
q1_input = Input(input_shape)
q2_input = Input(input_shape)
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#) 

```
model.add(Embedding(vocab_size, EMBEDDING_DIM, input_length=MAX_QUES_LEN,
                    embeddings_initializer=embeddings_init))
```

```
# 1st is a conv layer with max-pooling with 300 neurons with kernel size as 15
model.add(Conv1D(300, 15, activation='tanh')) # Conv1D
model.add(MaxPooling1D(pool_size=2))
```

```
# next will be deep network of 3 fully-connected layer with 50 neuron each

model.add(Dense(50, activation='tanh'))
model.add(Dense(50, activation='tanh'))
model.add(Dense(50, activation='tanh'))

# Generate the encodings (feature vectors) for the two images
model_q1 = model(q1_input)
model_q2 = model(q2_input)

# concatenating the outputs from the 2 model-layers of the siamese network
concat = keras.layers.Concatenate()([model_q1, model_q2])
flatcat = Flatten()(concat)

# Add a dense layer with sigmoid to generate the similarity score
prediction = Dense(1, activation='tanh')(flatcat)

# Connect the inputs with the outputs
siamese_net = models.Model(inputs=[q1_input, q2_input], outputs=prediction)

return siamese_net, model

# creating the network
siamese_net, seqmodel = get_siamese_model((MAX_QUES_LEN,))
siamese_net.summary()
seqmodel.summary()
```



Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)



WARNING: Logging before flag parsing goes to stderr.

W0805 03:39:46.018667 140377308571520 deprecation_wrapper.py:119] From /usr/local
 W0805 03:39:46.067380 140377308571520 deprecation_wrapper.py:119] From /usr/local
 W0805 03:39:46.261919 140377308571520 deprecation_wrapper.py:119] From /usr/local
 W0805 03:39:46.294217 140377308571520 deprecation_wrapper.py:119] From /usr/local
 2119

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 300)	0	
input_2 (InputLayer)	(None, 300)	0	
sequential_1 (Sequential)	(None, 143, 50)	6479450	input_1[0][0] input_2[0][0]
concatenate_1 (Concatenate)	(None, 143, 100)	0	sequential_1[1][sequential_1[2][
flatten_1 (Flatten)	(None, 14300)	0	concatenate_1[0]
dense_4 (Dense)	(None, 1)	14301	flatten_1[0][0]

Total params: 6,493,751

Trainable params: 6,493,751

Non-trainable params: 0

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 300, 300)	5109000
conv1d_1 (Conv1D)	(None, 286, 300)	1350300
max_pooling1d_1 (MaxPooling1	(None, 143, 300)	0
dense_1 (Dense)	(None, 143, 50)	15050
dense_2 (Dense)	(None, 143, 50)	2550
dense_3 (Dense)	(None, 143, 50)	2550

Total params: 6,479,450

Trainable params: 6,479,450

Non-trainable params: 0

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)



▼ FIT THE DATA

```
# Define optimiser, learning rate
optimizer = SGD(lr = 0.01)
siamese_net.compile(loss="mean_squared_error",optimizer=optimizer,metrics=['accuracy'])
history = siamese_net.fit([encoded_train_q1,encoded_train_q2], df_train_label, validation_da
```

```
W0805 03:39:50.902232 140377308571520 deprecation_wrapper.py:119] From /usr/local
```

```
W0805 03:39:51.193379 140377308571520 deprecation_wrapper.py:119] From /usr/local
```

```
W0805 03:39:51.259613 140377308571520 deprecation_wrapper.py:119] From /usr/local
```

```
Train on 16000 samples, validate on 4000 samples
```

```
Epoch 1/20
```

```
- 22s - loss: 0.2373 - acc: 0.5990 - val_loss: 0.2351 - val_acc: 0.5992
```

```
Epoch 2/20
```

```
- 15s - loss: 0.2222 - acc: 0.6422 - val_loss: 0.2305 - val_acc: 0.6170
```

```
Epoch 3/20
```

```
- 15s - loss: 0.2113 - acc: 0.6691 - val_loss: 0.2293 - val_acc: 0.6138
```

```
Epoch 4/20
```

```
- 16s - loss: 0.1998 - acc: 0.6974 - val_loss: 0.2265 - val_acc: 0.6338
```

```
Epoch 5/20
```

```
- 16s - loss: 0.1871 - acc: 0.7258 - val_loss: 0.2260 - val_acc: 0.6392
```

```
Epoch 6/20
```

```
- 16s - loss: 0.1748 - acc: 0.7508 - val_loss: 0.2304 - val_acc: 0.6395
```

```
Epoch 7/20
```

```
- 16s - loss: 0.1601 - acc: 0.7794 - val_loss: 0.2323 - val_acc: 0.6415
```

```
Epoch 8/20
```

```
- 16s - loss: 0.1432 - acc: 0.8158 - val_loss: 0.2405 - val_acc: 0.6325
```

```
Epoch 9/20
```

```
- 16s - loss: 0.1276 - acc: 0.8435 - val_loss: 0.2480 - val_acc: 0.6272
```

```
Epoch 10/20
```

```
- 16s - loss: 0.1109 - acc: 0.8740 - val_loss: 0.2521 - val_acc: 0.6225
```

```
Epoch 11/20
```

```
- 16s - loss: 0.0949 - acc: 0.8994 - val_loss: 0.2644 - val_acc: 0.6175
```

```
Epoch 12/20
```

```
- 16s - loss: 0.0819 - acc: 0.9193 - val_loss: 0.2650 - val_acc: 0.6148
```

```
Epoch 13/20
```

```
- 16s - loss: 0.0695 - acc: 0.9374 - val_loss: 0.2666 - val_acc: 0.5980
```

```
Epoch 14/20
```

```
- 16s - loss: 0.0574 - acc: 0.9546 - val_loss: 0.2720 - val_acc: 0.6038
```

```
Epoch 15/20
```

```
- 16s - loss: 0.0480 - acc: 0.9647 - val_loss: 0.2879 - val_acc: 0.5958
```

```
Epoch 16/20
```

```
- 16s - loss: 0.0402 - acc: 0.9734 - val_loss: 0.2891 - val_acc: 0.6012
```

```
Epoch 17/20
```

```
- 16s - loss: 0.0329 - acc: 0.9811 - val_loss: 0.3126 - val_acc: 0.5948
```

```
Epoch 18/20
```

```
- 16s - loss: 0.0277 - acc: 0.9857 - val_loss: 0.2973 - val_acc: 0.5975
```

```
Epoch 19/20
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)


```
print('Number of training samples = {}'.format(num_of_train))
print('Number of validation samples = {}'.format(num_of_valid))
print('Vocabulary size = {}'.format(vocab_size))
```



```
↳ Number of training samples = 8000  
    Number of validation samples = 2000  
    Vocabulary size = 17030
```

▼ PLOT GRAPHS


```
# Plotting training and validation loss (from lecture 2 notes)  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
epochs = range(1, len(loss)+1)  
  
plt.plot(epochs, loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



```
# Plotting training and validation accuracy (from lecture 2 notes)  
acc = history.history['acc']  
val_acc = history.history['val_acc']  
epochs = range(1, len(acc)+1)  
  
plt.plot(epochs, acc, 'bo', label='Training acc')  
plt.plot(epochs, val_acc, 'b', label='Validation acc')  
plt.title('Training and validation accuracy')  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.legend()
```

Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

↳



Automatic document saving has been pending for 26 minutes. Reloading may fix the problem. [Save and reload the page.](#)

